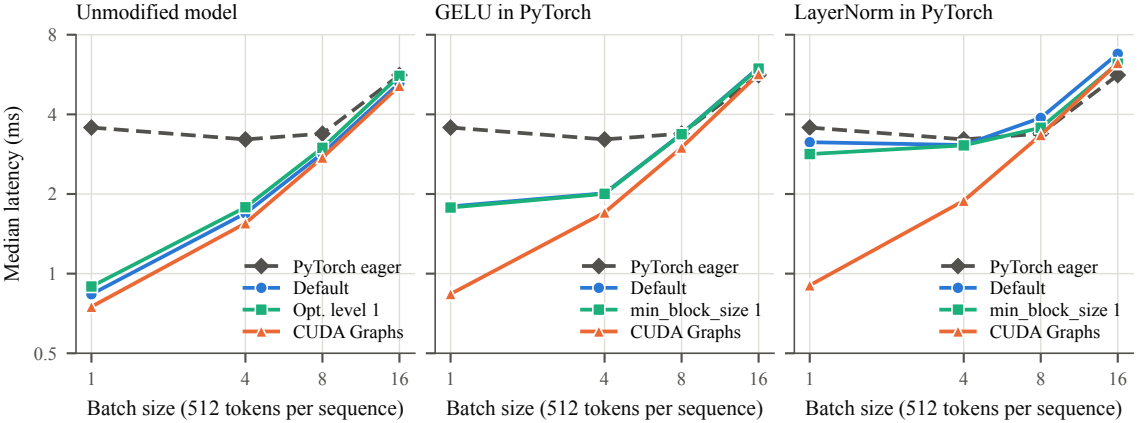
Unmodified model
GELU in PyTorch
LayerNorm in PyTorch
Median latency (ms)
8
4
2
1
0.5
PyTorch eager
Default
Opt. level 1
CUDA Graphs
min_block_size 1
1
4
8
16
Batch size (512 tokens per sequence)